

Claro

Banking support with verified outcomes

Immediate Spanish and Portuguese chat
with verified cases for human review

We keep organizer analysis, authored language scenarios and public demo fixtures
separate.

4.43M

transactions audited

42

distinct customer
utterances
in organizer transcripts

One conversation from inquiry to follow-up

Confirmar caso de revisión

Se creará un caso de revisión en este sandbox. No se realizará un reembolso, bloqueo de tarjeta ni movimiento de dinero.

Un cargo no reconocido requiere revisión humana. Con confirmación creamos un caso de prueba. No bloqueamos tarjetas ni prometemos reembolsos.

REPORTE QUE RECIBIRÁ EL ANALISTA

No reconozco el cargo de Tienda Demo; lo noté después de una llamada. Mi correo de prueba es [REDACTED].

Revisa el texto antes de confirmar. Los datos sensibles identificados se ocultan.

INFORMACIÓN PENDIENTE

- Se requiere valoración humana del reporte del cliente.

Cancelar✓ Confirmar creación

1 Start in chat

Automatic ES/PT language.
Clarify the transaction in chat.

2 Confirm and verify

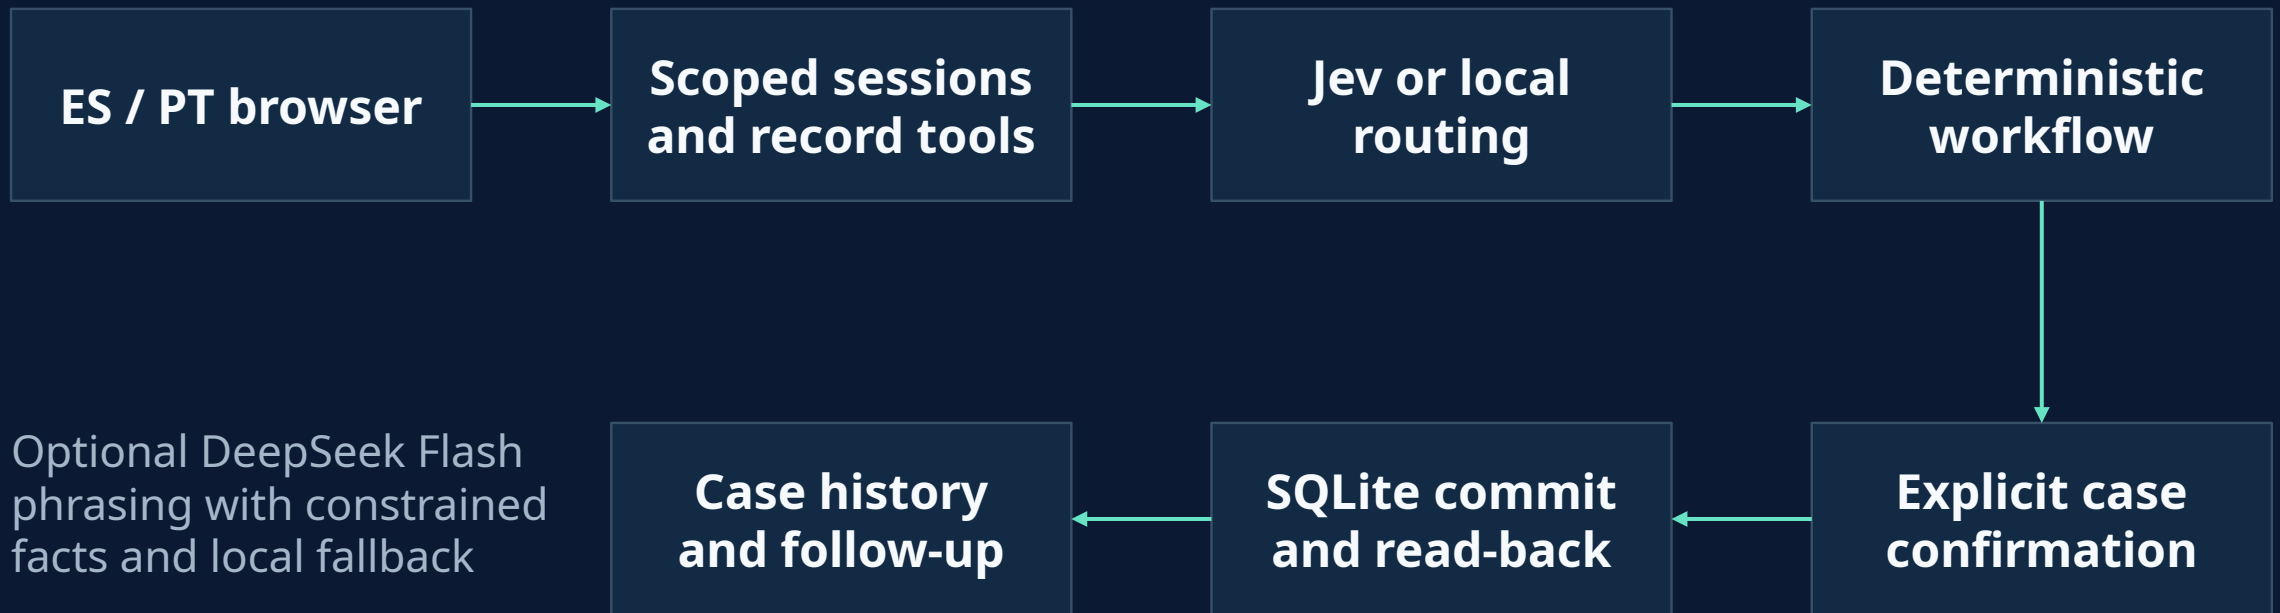
Commit one sandbox case.
Read it back before success.

3 Persist follow-up

Ask a specific question.
Reply in the same composer.

Permissions and actions stay in service code

Optional providers are implemented and mock-tested. Live inference remains unverified.



Opaque sessions + CSRF

Customer/workspace isolation

Idempotent confirmed writes

Fraud candidates did not earn release

4,425,008 validated transactions

Contracts, quarantine and source hashes
Strictly prior customer/currency history
Two full runs with identical feature bytes

Temporal evaluation

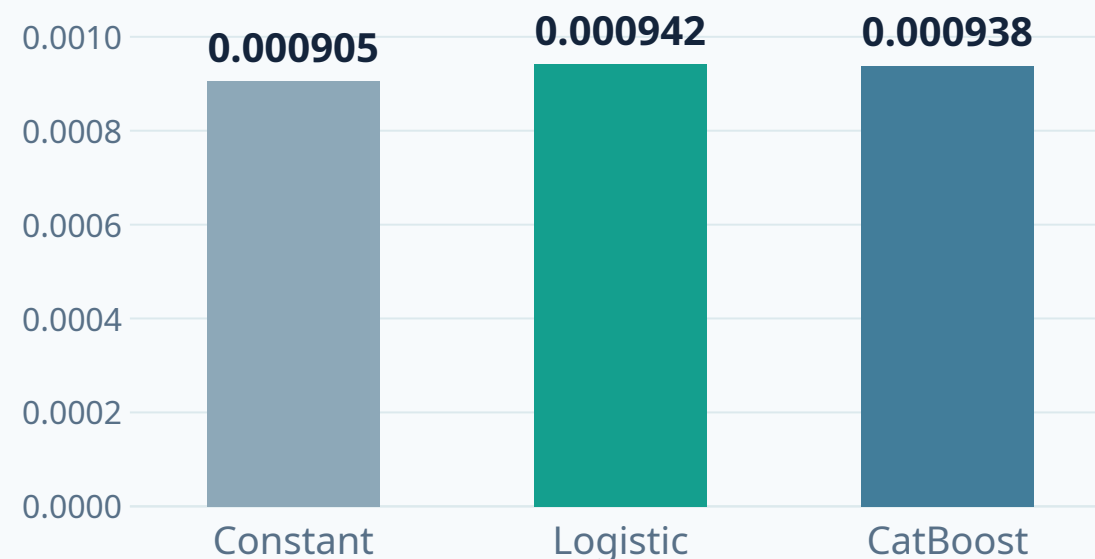
Train before January 2026
Calibrate January, select February
Test March through the June snapshot

Synthetic retrospective labels. Demo fixtures receive no fraud probability.

Claro / Factored 2026 / 4

451,823 test transactions / 409 fraud labels

Average precision



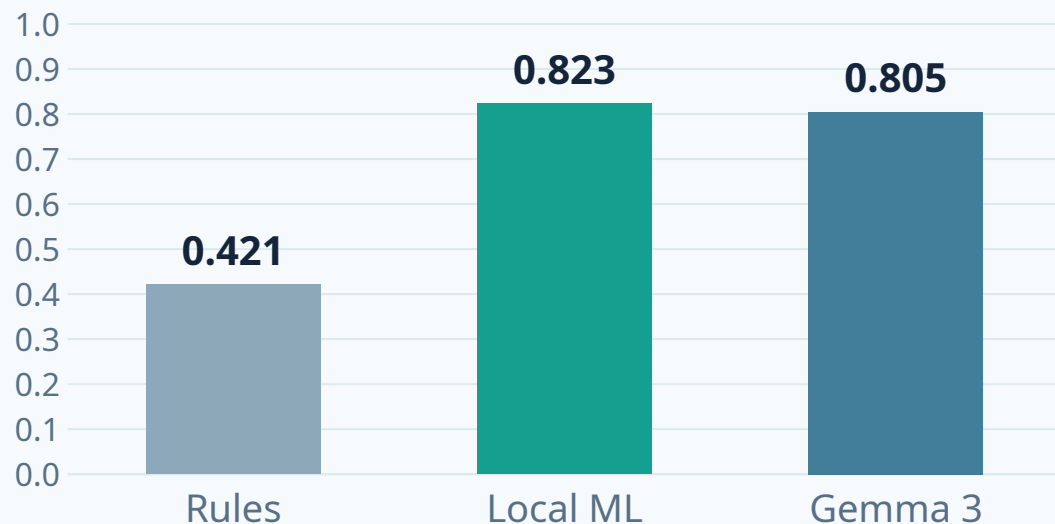
4 / 409

fraud labels reached at 1% review capacity by each learned candidate

Measured language and workflow results

Intent classification

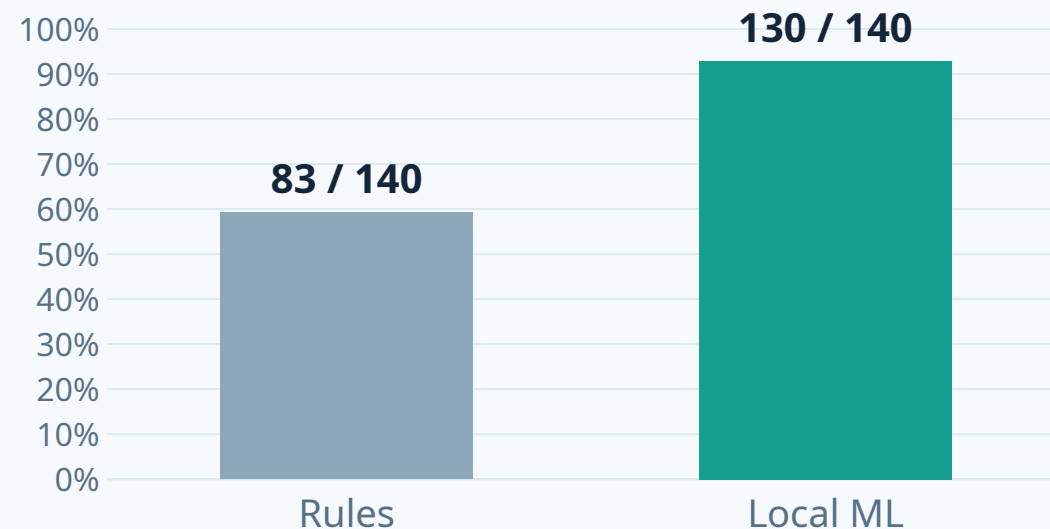
140 frozen authored cases / macro-F1



Historical intent labels: 117/140

Current chat regression

140 exposed cases / correct regression outcomes



60/60 handoffs; 10 workflow errors

Current six status/currency strata: 130/140 each. Same authored utterances.

Stronger case grounding scorer. Exposed authored regressions, no independent human review.

Demo and the route to a bank integration



Prior v1.1 hosted demo

<https://claro-banking-hackathon-2026.onrender.com>

Deliberate trade-offs

Single worker and SQLite
Simple deployment. Free-host restarts
can reset sandbox state.

Optional external models
Usage and tariff estimates are logged.
Live provider results are still unverified.

Before real customers

Bank identity, approved policies, durable
storage and independent ES/PT review.